

Computational Physics

PHYS 6260

Parallel Programming OpenMP

Announcements:

- HW6: Due Friday 2/28
- Project proposal due Friday 3/7

We will cover these topics

- Types of parallelism
- OpenMP

Lecture Outline

Optimization

- Getting performance out of your code means
 - Picking the right algorithm
 - Implementing the algorithm efficiently
- We previously covered about picking the most appropriate algorithm and saw some examples of speed-ups
- For performance in the implementation
 - We need to understand how the computer's CPU/GPU and architecture works
 - To exceed a single core's performance, we must go parallel

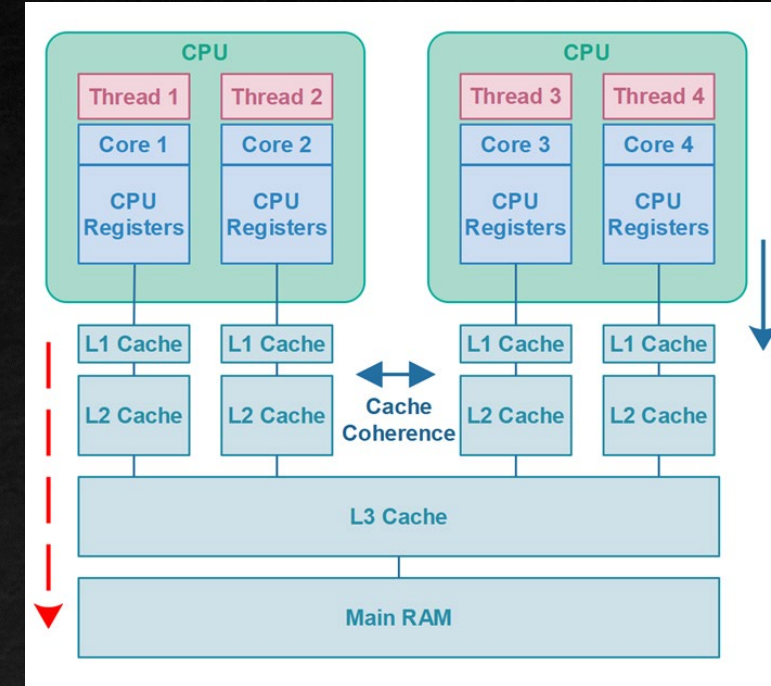
CPU + Memory System

- Memory hierarchy
 - Data is stored in main memory (RAM)
 - Below this, there are multiple levels of cache on the CPU (L3, L2, L1)
 - A line of memory is moved into cache – you amortize the costs if you put all the data in the line
 - Data is moved to the registers in the CPU – this is where the computation occurs
- It is expensive to move data from main memory to the registers
 - We need to exploit the cache
 - For arrays, loop over data such that you operate on elements that are adjacent in memory
 - In Python, you don't need to worry about this
 - But for compiled languages, you need to organize your loops optimally

CPU + Memory System


| Computer Action | Avg Latency | Normalized Human Time |
|---------------------------|---------------|-----------------------|
| 3GhzCPU Clock cycle 3Ghz | 0.3 ns | 1 s |
| Level 1 cache access | 0.9 ns | 3 s |
| Level 2 cache access | 2.8 ns | 9 s |
| Level 3 cache access | 12.9 ns | 43 s |
| RAM access | 70 - 100ns | 3.5 to 5.5 min |
| NVMe SSD I/O | 7-150 μ s | 2 hrs to 2 days |
| Rotational disk I/O | 1-10 ms | 11 days to 4 mos |
| Internet: SF to NYC | 40 ms | 1.2 years |
| Internet: SF to Australia | 183 ms | 6 years |
| OS virtualization reboot | 4 s | 127 years |
| Virtualization reboot | 40 s | 1200 years |
| Physical system reboot | 90 s | 3 Millenia |

Table 1: Computer Time in Human Termsⁱ




Arrays

- Row vs. Column major: $A(m,n)$
 - First index is called the row
 - Second index is called the column
 - Multi-dimensional arrays are flattened into one-dimensional sequence for storage
- Row-major (C, python): rows are stored one after the other
- Column-major (Fortran, matlab): columns are stored one after the other
- Ordering matters for:
 - Passing arrays between languages
 - Deciding which index to loop over first



Row major



Column major


Arrays

- This is why in C, you'd want to loop as


```
double A[M][N];  
for (i = 0; i < M; i++) {  
    for (j = 0; j < N; j++) {  
        A[i][j] = ...;  
    }  
}
```

- And in Fortran / matlab

```
double precision :: A(M,N)  
do j = 1, N  
    do i = 1, M  
        A(i,j) = ...  
    enddo  
enddo
```



Row major



Column major

Types of parallelism

- Flynn's taxonomy classifies computer architectures
- 4 classifications: single or multiple data; single or multiple instruction
- **Single instruction, single data (SISD)**
 - Typical application on your computer – no parallelism
- **Single instruction, multiple data (SIMD)**
 - The same instruction set is done to multiple pieces of data all at once
 - Compile-time vectorization [optimization](#); GPUs
- **Multiple instructions, single data (MISD)**
 - Not very interesting or useful
- **Multiple instructions, multiple data (MIMD)**
 - This is what we classify as parallel computing

Types of parallelism




Figure 2. Intel® AVX and Intel® SSE data types




Figure 3. SIMD versus scalar operations

Trivially parallel

- Sometimes our tasks are trivially parallel
 - No communication is needed between processes
- Examples: ray tracing, Monte Carlo, parameter sweeps
 - Each realization can do its work independently
 - At the end, we may need to do some simple processing of all of the results
- Large data analysis
 - Imagine that you have many datasets and a reduction pipeline to work on them
 - Use multiple processors to work on the different data files as resources become available
 - Each file is processed on a single core

Trivially parallel with shell scripts

- Example: data analysis – launch independent jobs
- This can be done through a shell script – no libraries necessary
- Loop over files
 - Run jobs until all of the processors are full
 - Use lock files to indicate a job is running
 - When resources become free, start up the next job
 - Let's look at some code
- Also see GNU parallel – very useful to run multiple commands at once
- Example: sorting all files named out*list by the 8th column in reverse order


```
$ ls out*list | parallel sort -gr -k 8 -o sorted_{} {}
```


How do we make our code parallel?

- There is no simple compiler flag “—make-this-parallel”
- We have to understand the algorithm and determine what parts are amenable to parallelism
- However, if the bulk of your work is in one specific piece (say, solving a linear system), you may get all that you need by using a library that is already parallel
 - Example: [hypre](#) for solving linear systems; [fftw](#) for FFTs
 - This will require minimal changes to your code

Shared Memory / OpenMP

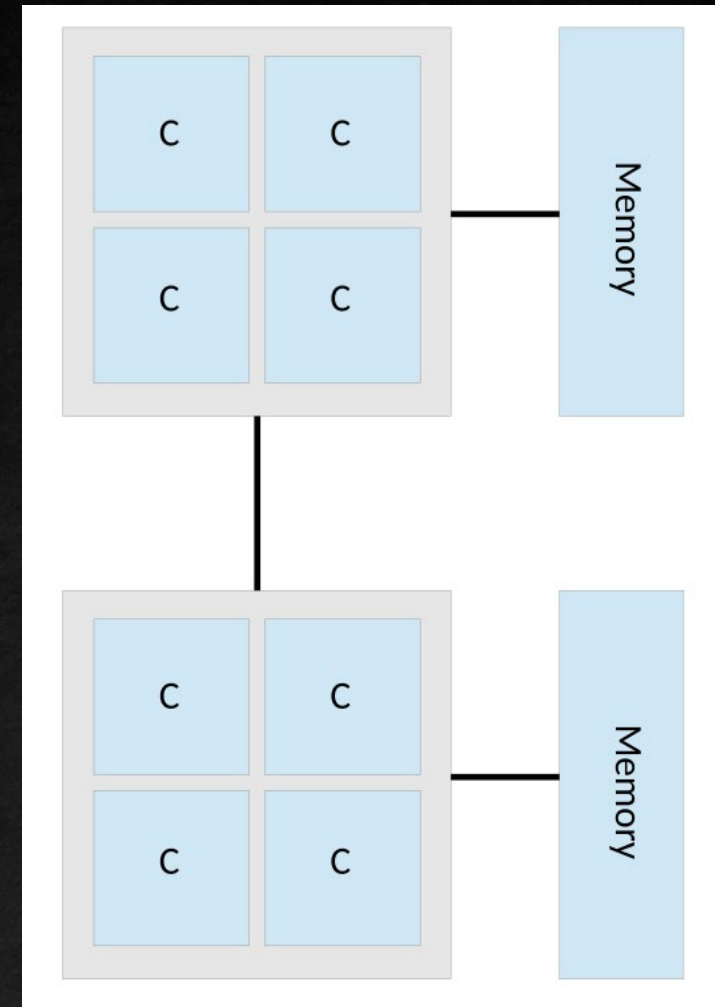
- Nodes consist of one or more multi-core (2-128) CPUs
- Everything can access the same pool of memory



Single 4-core chip and its pool of memory

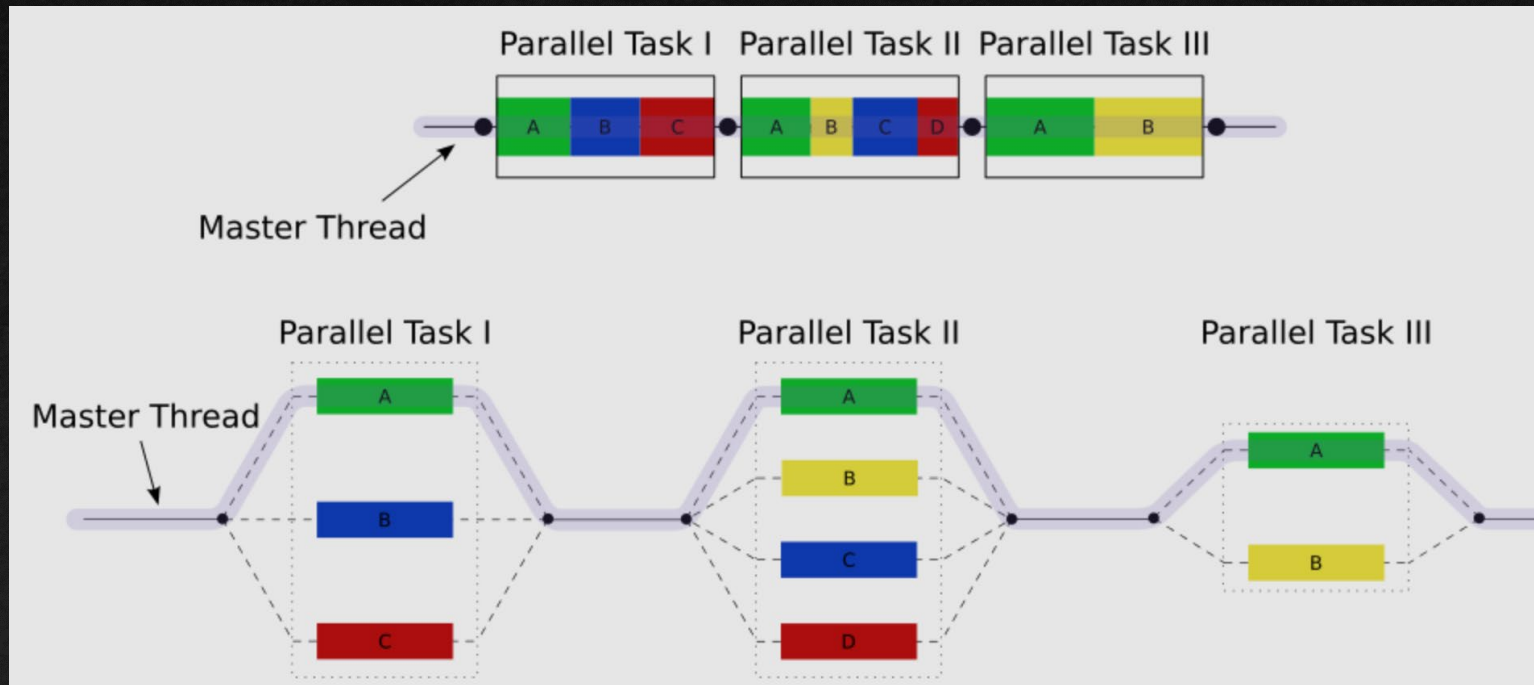
Shared Memory / OpenMP

- Some machines are more complex – multiple chips each with their own pool of local memory
- Latency may be higher when going “off-chip”
- Best performance will require you to know the machine’s architecture and optimizing the run-time parameters for it



Shared Memory / OpenMP

- When using OpenMP, threads are spawned as needed
- When you run the program, there is one thread – the main thread
- When you enter a parallel region, multiple threads run concurrently



Shared Memory / OpenMP

- OpenMP is initiated through directives or pragmas
- Look like comments unless you tell the compiler to interpret them
- Environment variable OMP_NUM_THREADS sets the number of threads
- Support for C, C++, Fortran
- Hello world

```
#include <stdio.h>
void main() {
    #pragma omp parallel
    printf("Hello world\n");
}
```

- Compile with `gcc -o hello -fopenmp hello.c`

Shared Memory / OpenMP

- In addition to using pragmas, there are a few functions that OpenMP provides to get the number of threads, the current thread, etc.

```
#include <stdio.h>
#ifdef _OPENMP
#include <omp.h>
#endif

void main() {

    int nthreads, thread_num;

#ifdef _OPENMP
    nthreads = omp_get_num_threads();
    thread_num = omp_get_thread_num();
    printf("[thread %d] Outside of parallel region = %d\n", thread_num, nthreads);
#pragma omp parallel
    {
        nthreads = omp_get_num_threads();
        thread_num = omp_get_thread_num();
        printf("[thread %d] Inside of parallel region = %d\n", thread_num, nthreads);
    }
#else
    printf("OpenMP not enabled.\n");
#endif
}
```


Shared Memory / OpenMP

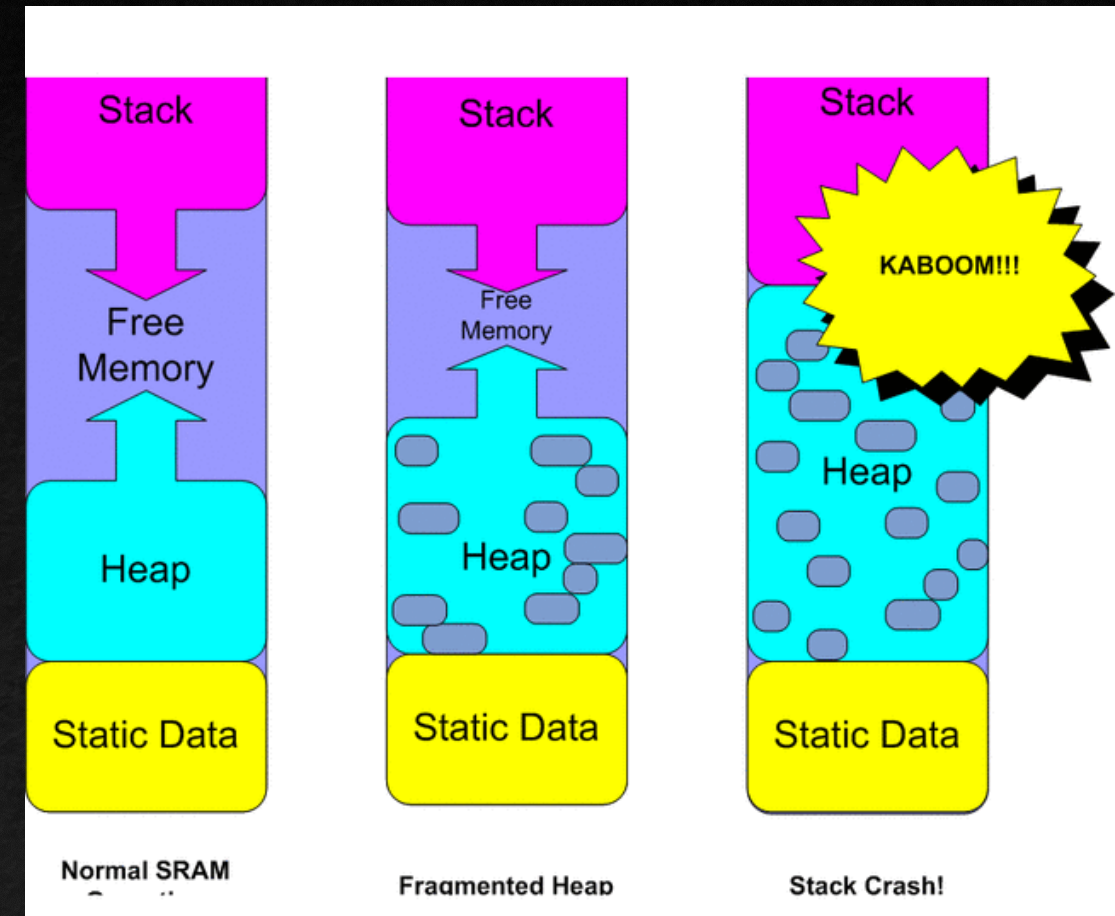
- Most (all?) modern compilers support OpenMP
- However, the performance across them can vary
- GCC does a reasonable job. Intel compilers are the fastest (this also includes SIMD vectorization)
- There is an overhead associated with spawning threads
- You will have to experiment to see if parallelism is worth it
- Some regions of code may not have enough work to offset the overhead

Shared Memory / OpenMP

- There will be a systemwide default for OMP_NUM_THREADS
- Things will still run if you use more threads than cores available
- But don't do this!
- This is called oversubscription, and will cause work to queue up
- It's better to use a maximum of 1 OpenMP thread per core
- Scaling: if you double the number of cores, does the code take $\frac{1}{2}$ the time?

Aside about memory: Stack vs. Heap

- Memory allocated at compile time is put on the stack, e.g. `double a[1000];`
- Stack memory has a fixed and somewhat small size
 - Managed by the operating system
 - You don't need to clean up / free this memory
- Dynamic allocation puts the memory on the heap
 - Much bigger pool
 - You are responsible for deallocating / freeing the memory
 - Otherwise, you will have memory leaks and the memory footprint will continuously increase



Shared Memory / OpenMP: Loops

- Parallel example of matrix multiplication
- Notice the private() calls
- Inside the loop, all threads will have access to all of the variables declared in the main program
- For some variables, we want a private copy on each thread, like the loop counters
 - These are put in the private() clause

```
#pragma omp parallel private(i,j,n)
{
    n = 0;
    #pragma omp for
    for (j = 0; j < N; j++) {
        for (i = 0; i < N; i++, n++) {
            a[n] = i + j;
        }
        x[j] = j;
        b[j] = 0.0;
    }
}
```

```
// Multiply
n = 0;
#pragma omp for
for (j = 0; j < N; j++) {
    for (i = 0; i < N; i++, n++) {
        b[i] = b[i] + a[n] * x[j];
    }
}
```


Shared Memory / OpenMP: Reduction

- Suppose that you are finding the minimum value of something or summing
- The loop over the data are spread across threads
- How do we get the data from each thread back to a single variable that all threads see?
- `reduction()` clause
 - Has both shared and private behaviors
 - Compiler ensures that the data is synchronized at the end

Shared Memory / OpenMP: Reduction

```
sum = 0;  
#pragma omp parallel for reduction(+:sa, sb) reduction(max:m)  
for (i = 0; i < N; i++) {  
    sa += A[i];  
    sb += B[i];  
    if (A[i]>m) m = A[i];  
    m = B[i]>m?B[i]:m;    // same result as: if (B[i]>m) m = B[i];  
}
```

sa will be the sum over array A

sb will be the sum over array B

m will be the max over both arrays

OpenMP Example: Relaxation

- Remember in the relaxation method to solve an elliptic PDE, we have

$$\phi_{i,j} = \frac{1}{4}(\phi_{i+1,j} + \phi_{i-1,j} + \phi_{i,j+1} + \phi_{i,j-1} - (\Delta x)^2 f_{i,j})$$

- Red-black Gauss-Seidel:
 - Update in place
 - First update the red cells with the black cells unchanged
 - Then update the black cells with the red cells unchanged
- Let's inspect a relaxation solver (in C). Remember that OpenMP is only supported for compiled programs.

