

Computational Physics

PHYS 6260

Parallel Programming OpenMP / MPI

Announcements:

- HW6: Due Friday 3/1
- Project proposal, Due Friday 3/7

We will cover these topics

- Finishing up OpenMP
- MPI

Lecture Outline

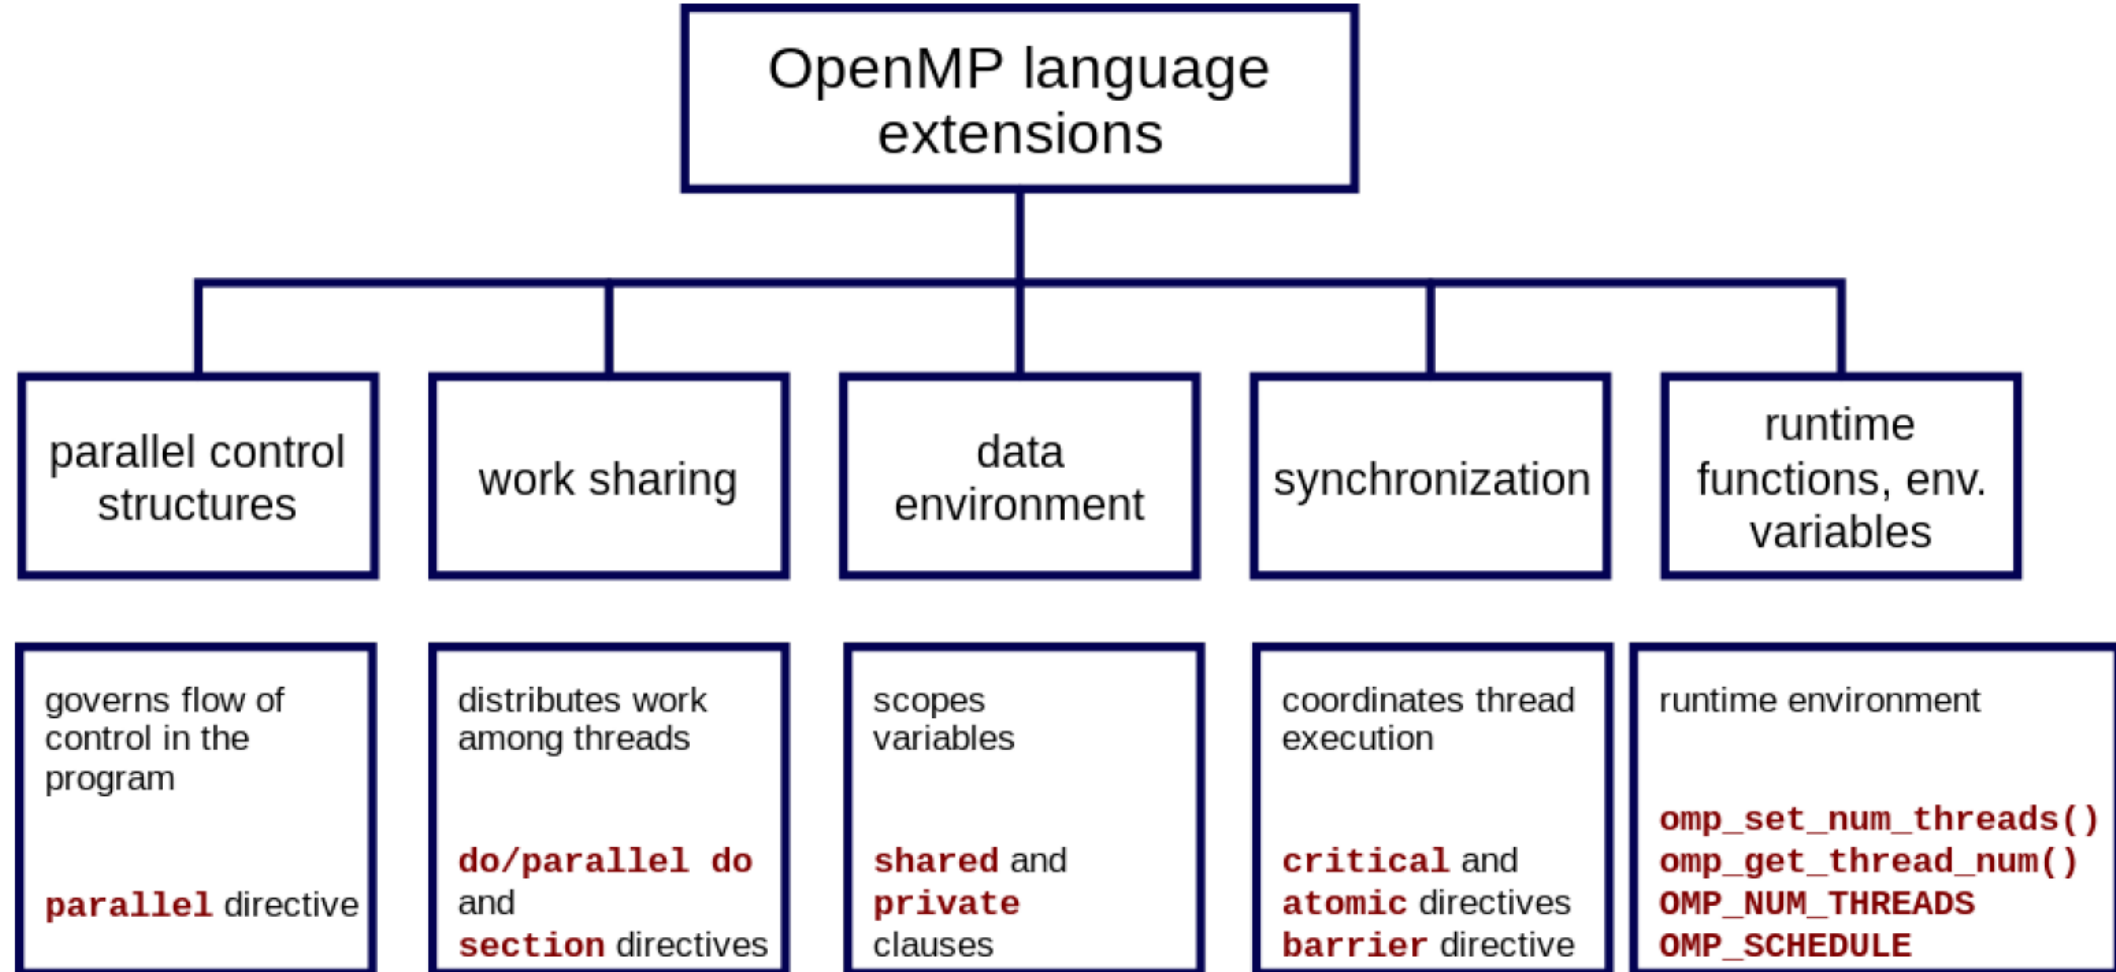
OpenMP: Critical sections

- Within a parallel region, sometimes you need to ensure that only one thread at a time can write to a variable
- Consider the following

```
if ( a[i][j] > maxa ) {  
    maxa = a[i][j]  
    imax = i  
    jmax = j  
}
```

- If this is in the middle of a loop, what happens if 2 different threads meet this criteria?
- Marking this section as critical will ensure only one thread changes variables at a time
- Warning: critical sections can be very slow

OpenMP: available modes



OpenMP: Porting to OpenMP

- You can parallelize your code piece-by-piece
- Since OpenMP directives look like comments to the compiler, your old version is still there
- Generally, you are not changing any of the original code – just adding directives

OpenMP: Advanced options

- “if” clause tells OpenMP only to parallelize a region if certain conditions are met (e.g. a test of the size of an array)
- “firstprivate” is like “private” except each copy is initialized to the value from the original value
- “schedule” affects the balance of the work distributed to the threads. Options: static, dynamic, guided, chunk sizes, number of threads

OpenMP: OpenMP in Python

- Python enforces a “global interpreter lock” that means only one thread can talk to the interpreter at any one time
 - OpenMP within pure python is not possible
- However, C extensions or Cython code called from python can use shared-memory parallelism
 - That is, the underlying code can have OpenMP directives

MPI Parallelism

MPI: Distributed parallelism

- The Message Passing Library (MPI) is the standard library for distributed parallel computing
- Now each core cannot directly see each other's memory
- You need to manage how the work is divided and explicitly send and receive messages between MPI processes

MPI: Hello world

- No longer do we simply use comments. We need to call routines in the MPI library.

```
MPI_Init(&argc, &argv);
MPI_Comm_size(MPI_COMM_WORLD, &size);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);

hostname[MAX_LENGTH-1] = '\0';
gethostname(hostname, MAX_LENGTH-1);
printf("P%04d/%04d: Hello world from %s\n",
        rank, size, hostname);

MPI_Finalize();
```


MPI: Hello world

- MPI jobs are run with a command line executable, usually called `mpirun`, where you specify the number of cores

```
mpirun -n 4 ./hello_mpi
```

- You need to install the MPI libraries on your machine
 - OpenMPI is the most popular on local machines
e.g. “`apt install openmpi-bin libopenmpi-dev`” on Debian systems
 - MVAPICH2 is the most common on HPC platforms

MPI: Concepts

- A separate instance of your program is run on each core – these are MPI processes
- Thread safety is not an issue here because each instance of the program is isolated from the others
- You need to tell the library the datatype of the variable you are communicating and its size (i.e. the buffer)
- Processors can be grouped together
 - Communicators label different groups
 - `MPI_COMM_WORLD` is the default communicator (all processes)

MPI: Concepts

- Many types of operations: send/receive, collective (e.g. broadcast, gather, scatter)
- There are 100+ functions, but one can do any message passing with only 6
 - MPI_Init
 - MPI_Comm_Size
 - MPI_Comm_Rank
 - MPI_Send
 - MPI_Recv
 - MPI_Finalize
- More efficient communication can be done with the more advanced functions
- System vendors will usually provide their own MPI implementation that is well matched with the hardware

MPI Example: Computing pi

- Let's calculate pi by computing the integral (as compared to the trivially parallel Monte Carlo method)

$$\int_0^1 \frac{1}{1+x^2} dx = \arctan(x) \Big|_0^1 = \frac{\pi}{4}$$


- We will divide the $[0,1]$ interval into N_{proc} pieces so each process only sees its own interval
- Each core computes the sum for its interval
- The “root” process will add all the integrals at the end
- Let's look at the code

MPI Example: Send / receive

- The main idea in MPI is sending messages between processes
- MPI_Send() and MPI_Recv() pairs provide this functionality
- This is a “blocking” send / receive
 - For sending code, the program resumes when it is safe to reuse the buffer
 - For receiving code, the program resumes when the message is received
- May cause network contention if the destination process is busy doing its own communication
- There are non-blocking sends, where you explicitly attach a buffer and the program continues after sending the data

MPI Example: Send / receive

- Simple example (mimics ghost zone filling)
- One each process, allocate an integer array of 5 elements
- Fill the middle 3 with a sequence (e.g. proc 0: [0, 1, 2], proc 1: [3, 4, 5], ...)
- Send messages to fill the left and right element with the corresponding element from the neighboring processes



MPI Example: Send / receive


- Good communication performance often requires staggering the communication
- This is accomplished through overlapping communication and computation
- For example, you would send the data, compute something, and then receive any data necessary
- A combined `MPI_Sendrecv()` call can help avoid deadlocking
- Let's look at the same code with such a call

MPI Example: Relaxation

- Let's do the same relaxation problem but now using MPI instead of OpenMP
- In the OpenMP version, we allocated a single array covering the entire domain
- All cores had access to the whole array
- In the MPI version, each processor will allocate a smaller array, covering only a portion of the entire domain
- Each process can only access their own portion
- Necessary when data exceeds the memory limits of a single node / computer


MPI Example: Relaxation

- We will decompose the domain in 1D
- Each process allocates a slab that covers the full y-extent of the domain
- Width in x-direction = n_x/n_{procs}
 - If not evenly divisible, some slabs will have differing widths
- Perimeter of 1 ghost zone surrounding each subdomain
- We will refer to a global index space $[0:n_x-1] \times [0:n_y-1]$
- Memory needs to spread across all cores




MPI Example: Relaxation

- Left set of ghost zones are filled by receiving a message from the core (slab) to the left





MPI Example: Relaxation

- Left set of ghost zones are filled by receiving a message from the core (slab) to the left
- Right set of ghost zones are filled by receiving a message from the process (slab) to the right
- Top and bottom ghost zones are physical boundaries



MPI Example: Relaxation

- Generally speaking, one wants to minimize the surface-to-volume ratio
- Reduces communication



MPI Example: Relaxation

- Most of the parallelism comes in the ghost zone filling
 - Fill left GZs by receiving data from the process to the left
 - Fill right GZs by receiving data from the process to the right
 - Send / receive pairs: we want to try to avoid contention (this can be very tricky, and people spend a lot of time worrying about this!)
- On the physical boundaries, we simply fill as usual
- The way we have written the code, our relaxation routine doesn't need to do any parallelism itself.
 - It just operates on the domain it is given
- For computing a norm, we will need to reduce the local sums across processes
- Let's look at the code

Weak versus Strong Scaling

- In assessing the parallel performance of code, there are two common methods
- **Strong scaling**: keep the problem size fixed and increase the number of cores
 - Eventually you will become work-starved
 - Scaling will stop when communication and overhead dominate
- **Weak scaling**: increase the amount of work in proportion to the number of cores
 - Perfect scaling will result in the same “wall clock” time for all core counts

Parallel debugging

- There are parallel debuggers, but they are usually expensive
- It's possible to spawn multiple gdb (GNU debugger) sessions, but this gets out of hand quickly – “`mpirun -n 2 xterm -e gdb ./a.out`”
- Print is still your friend
 - Run a small of a problem as possible on as few cores as necessary
- Some round off differences are to be expected from sums (different order of operation ... numeric addition is not associative)

Hybrid parallelism

- To get good performance on current HPC platforms (>1k cores), you need to use hybrid parallelism
- OpenMP within a node / socket, MPI across nodes
- For example, in our MPI relaxation code, we could split the loops over each subdomain over multiple cores on a node using OpenMP
- Then we have MPI to communicate across nodes and OpenMP within the nodes
- The hybrid approach is often needed to get the best performance on big machines