

Computational Physics

PHYS 6260

Machine Learning: Neural Networks II

Announcements:

- Spring Break next week!
- Project progress report: Due Friday 3/28


We will cover these topics

- Constructing a NN from scratch
- Hidden layers
- Image classification

Lecture Outline


Neural network overview

- Neural networks are divided into **layers**
 - There's always an input layer – it doesn't do any processing – just accepts the input
 - There is always output layers
- Within a layer, there are neurons or **nodes**
 - For input, there will be one node for each input variable
- Every node in the first layer connects to every node in the next layer
 - The **weight** associated with the **connection** can vary – these are the matrix elements




Hidden layers

- We can add more parameters by adding another layer of nodes
- Layers between the input and output are known as **hidden layers**
- They introduce non-linear combinations of the inputs into the model



Hidden layers

- For a hidden layer of dimension k :
 - Inputs: $x \in \mathbb{R}^n$
 - Outputs: $z \in \mathbb{R}^m$
 - A is an $(m \times k)$ matrix
 - B is a $(k \times n)$ matrix
 - The product AB is $(m \times n)$, as we had before without a hidden layer
- **Universal approximation theorem**: single layer networks can represent any continuous function
- From now on, we will not use an α , so the sigmoid functions are the same in each layer




Hidden layers

- We transform the input in two steps:

$$\tilde{z} = g(Bx)$$

$$z = g(A\tilde{z})$$




Hidden layers

- Minimize to find the A, B matrices

$$f(A_{ls}, B_{ij}) = \sum_{l=1}^m (z_l - y_l)^2$$

$$\tilde{z}_i = g\left(\sum_{j=1}^n B_{ij}x_j\right)$$

$$z_l = g\left(\sum_{s=1}^k A_{ls}\tilde{z}_s\right)$$



Hidden layers

- Minimize to find the A, B matrices

$$f(A_{ls}, B_{ij}) = \sum_{l=1}^m (z_l - y_l)^2$$


$$\tilde{z}_i = g \left(\sum_{j=1}^n B_{ij} x_j \right)$$

$$z_l = g \left(\sum_{s=1}^k A_{ls} \tilde{z}_s \right)$$

- We need to minimize both sets of weights (A and B matrices)
- In practice, we do them one at a time, with each seeing the result from its layer
- This process is called **backpropagation**

Backpropagation

- Backpropagation in NNs use the errors at the end to change the weights that came in earlier in the network
- In the evaluation step, we progress through the NN in a forward direction (input → hidden layers → output)
- Backpropagation is the process of taking the errors that we compute at the output layer and moving them backwards to the hidden layer




Backpropagation: gradient descent

- Perform gradient descent on **A** and **B** separately
- This is the strength of backpropagation and descent versus some “canned” minimization routine – we are not optimizing the entire system
- Differentiating our error \vec{e} through numerous chain rules gives:

$$\Delta A = -2\eta \vec{e} \circ \vec{z} \circ (1 - \vec{z}) \cdot \tilde{z}^T$$
$$\Delta B = -2\eta \tilde{e} \circ \tilde{z} \circ (1 - \tilde{z}) \cdot \vec{x}^T$$


where

$$\tilde{e} = A^T \vec{e} \circ \vec{z} \circ (1 - \vec{z}) \approx A^T \vec{e}$$



Hidden layers

- Usually only a single hidden layer is needed
- In general, you want fewer nodes in your hidden layer(s) than in your input layer
 - Reasonable choice: $(n) \text{ inputs} > (k) \text{ hidden nodes} > (m) \text{ outputs}$
- Interactive hidden layers: <https://playground.tensorflow.org/>




Another example: signal analysis (signal_test_m4)

- Consider a noisy signal that we expect to lie in 1 of 4 frequency bands: $f = \{1, 2, 3, 4\}$
- The clean signal should be: $s = \cos(2\pi ft)$
- We are given n points in the form
$$x_i = \cos(2\pi ft) + 5r_i$$
- $r_i \in [-1, +1]$ is a random number, making the signal-to-noise ratio small
- We will have 4 outputs with a 1 in the position corresponding to the frequency, e.g. 1 Hz: $[1, 0, 0, 0]$; 2 Hz: $[0, 1, 0, 0]$
- We'll train a NN on known input/output pairs and then test with unknown pairs – **can we recover the frequency?**


Another example: signal analysis

- Here's a single frequency ($f=1$) sample data set
- See the very low SNR
- We'll use 5 epochs to train the NN
- Learning rate $\eta = 0.05$




Another example: signal analysis

- $k = 2$ hidden layers
- Here is a histogram of Δf
 - Zero meaning that we predicted the frequency correctly
- 1000 random datasets in the training set



Another example: signal analysis


- $k = 2$ hidden layers
- Here is a histogram of Δf
 - Zero meaning that we predicted the frequency correctly
- Now on the testing datasets with 100 sets
- Let's try to increase the hidden layers




Another example: signal analysis

- $k = 4$ hidden layers

Training set




Random data




Another example: signal analysis

- $k = 8$ hidden layers

Training set




Random data




Another example: signal analysis

- $k = 32$ hidden layers. Much better!

Training set




Random data



Another example: signal analysis

- Is a NN the best choice for this problem?
 - ML isn't always the answer
- We could imagine doing this sample example with an FFT
- Take the FFT of the test signal and return the frequency with the maximum power



Another example: signal analysis (`fft_compare.py`)

- Perform this same FFT analysis 100 times
- It only has a 50% success rate
- But the error is not dominated by the adjacent frequency ($\Delta f = 1$)




Image classification

- We'll try to recognize a digit (0—9) from an image of a handwritten digit
- Use the MNIST dataset (see Canvas since the original location is password protected now)
 - Popular dataset for testing ML techniques
 - Training (testing) set is 60k (10k) images from ~250 different people
 - Correct answer is known for both sets so we can test our performance
- Image details: 28x28 pixels, grayscale (0—255 intensity)
- The best ML algorithms can get accuracy >99%

Image classification: NN details

- Input layer will have 784 nodes (pixels)
- Output layer will have 10 nodes
 - Array with an entry for each possible digit
 - E.g. "3" would be represented as [0, 0, 0, 1, 0, 0, 0, 0, 0, 0]
 - We'll start with a hidden layer size of 100
- Train on all 60k training datasets
- Test on all 10k testing datasets




Image classification: NN details

- Let's look through the code `19_char_recognition.py`
- With the default configuration in the code, we achieve 95-96% accuracy!

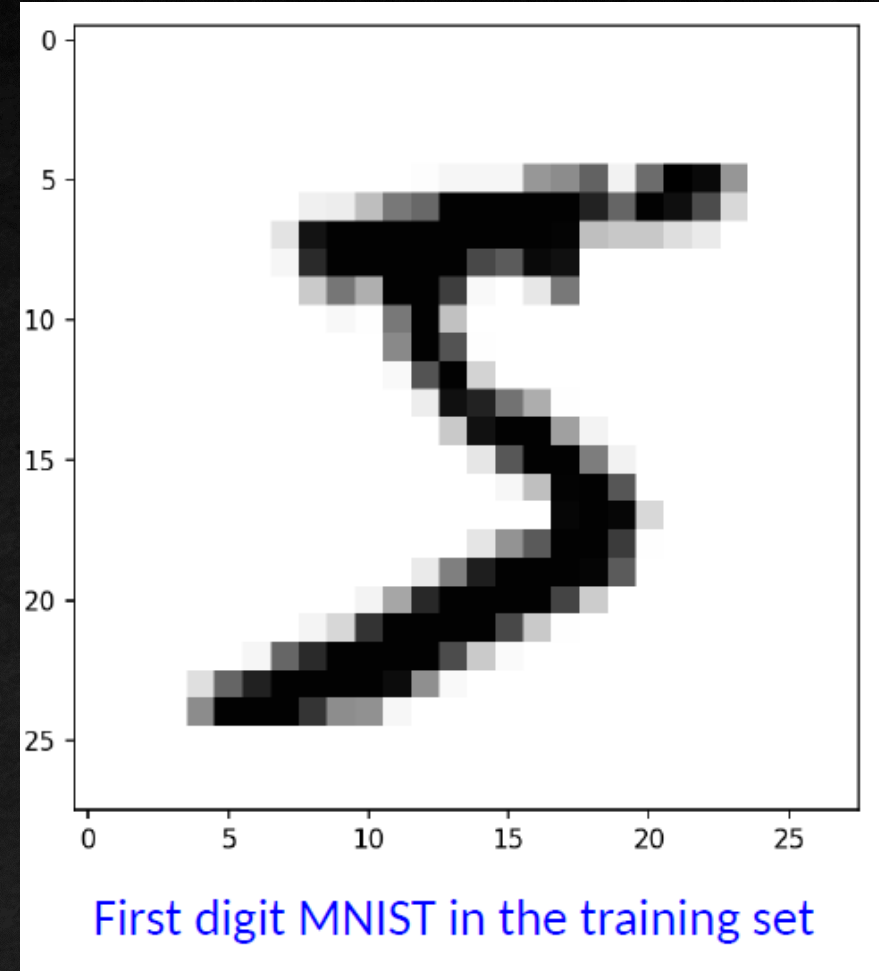


Image classification: some failures





Image classification: some failures





Image classification: some failures





Image classification: image classification matrices

- The weights (A and B matrices) seem symmetric around 0
- Interestingly with more training, the width of the distribution seems to grow




Image classification: Number of epochs

- When we use the full training set (60k images), the number of epochs (passes through the training data) doesn't matter much

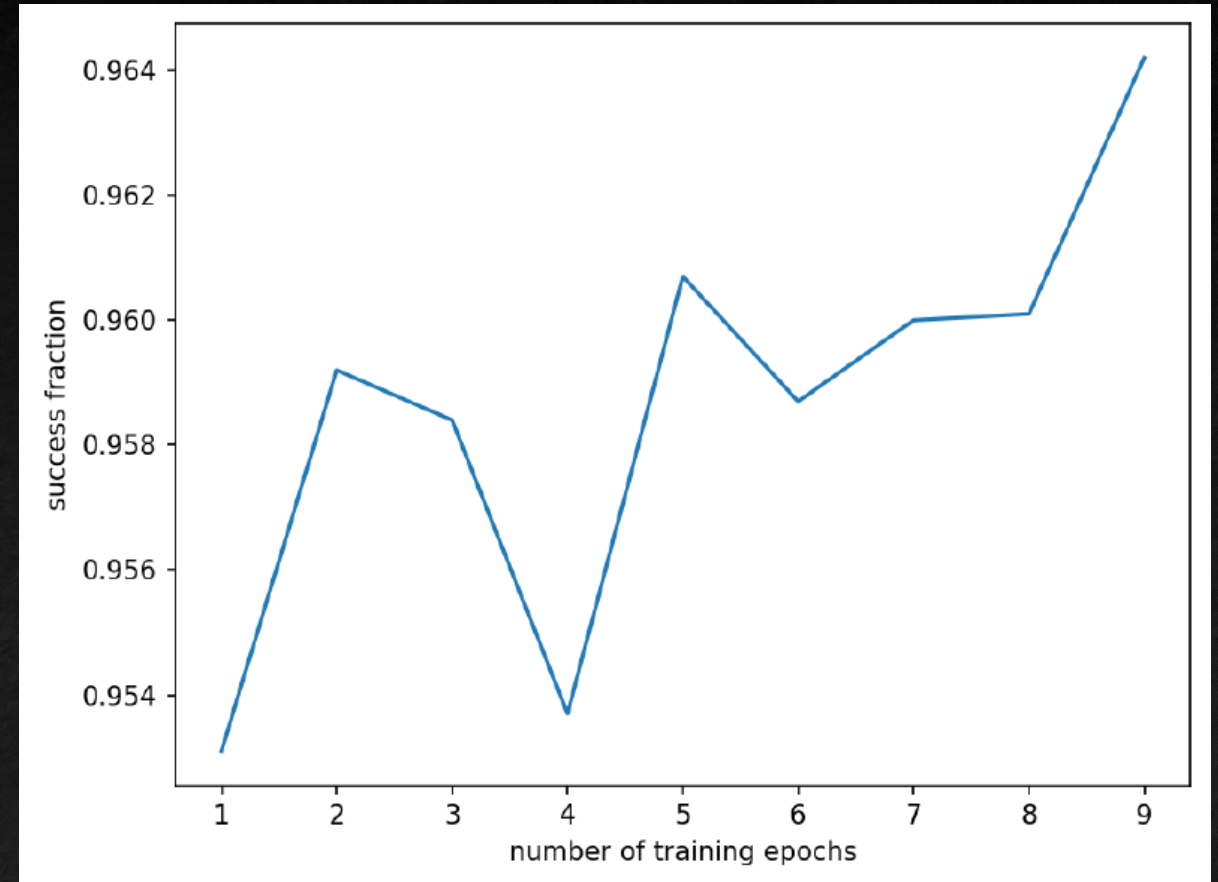


Image classification: Number of epochs

- No surprise: the larger the training set, the better the NN performs





Image classification: Hidden nodes / layer

- The larger the hidden layer, the better the performance



Deep Learning (after Spring Break)

- A deep neural network is one with many hidden layers (certainly >1 hidden)
- Surprisingly nice discussion on [StackExchange](#)
- [Free textbook](#)
- Another [free textbook](#)
- There are other ML algorithms aside from neural networks
 - [2019 review](#) of ML and physical sciences